

Embedded Systems Design

Lecture 9

Yongsoo Joo

VM for Embedded Systems

- Virtual memory
 - Not an issue for embedded systems
 - A single dedicated application without OS support
 - Things changed
 - Many of modern embedded systems rely on OS

Memory Management

- Goals


- Convenient abstraction for programming
- Provide isolation between different processes
- Allocate scarce physical memory resources across processes
 - Especially important when memory is heavily contended for

- Mechanisms

- Virtual address translation
- Paging and TLBs
- Page table management

- Policies

- Page replacement policies



Programmer's view of memory

Virtual Memory

- What is VM?
 - The basic abstraction provided by the OS for memory management
 - Enables programs to be executed without requiring their entire address space to be resident in physical memory
 - Called “Demand paging”
- Observation
 - Many programs don’t use all of their code or data
 - Ex) Branches never taken, variables never accessed, etc.
 - No need to allocate memory for it until it’s used

Virtual Memory

- Functions

- OS uses VM to adjust amount **of physical memory** allocated **to each process** based on its run-time behavior
- VM also provides isolation among processes
- Each process has its own isolated address space
- One process cannot access memory addresses in others

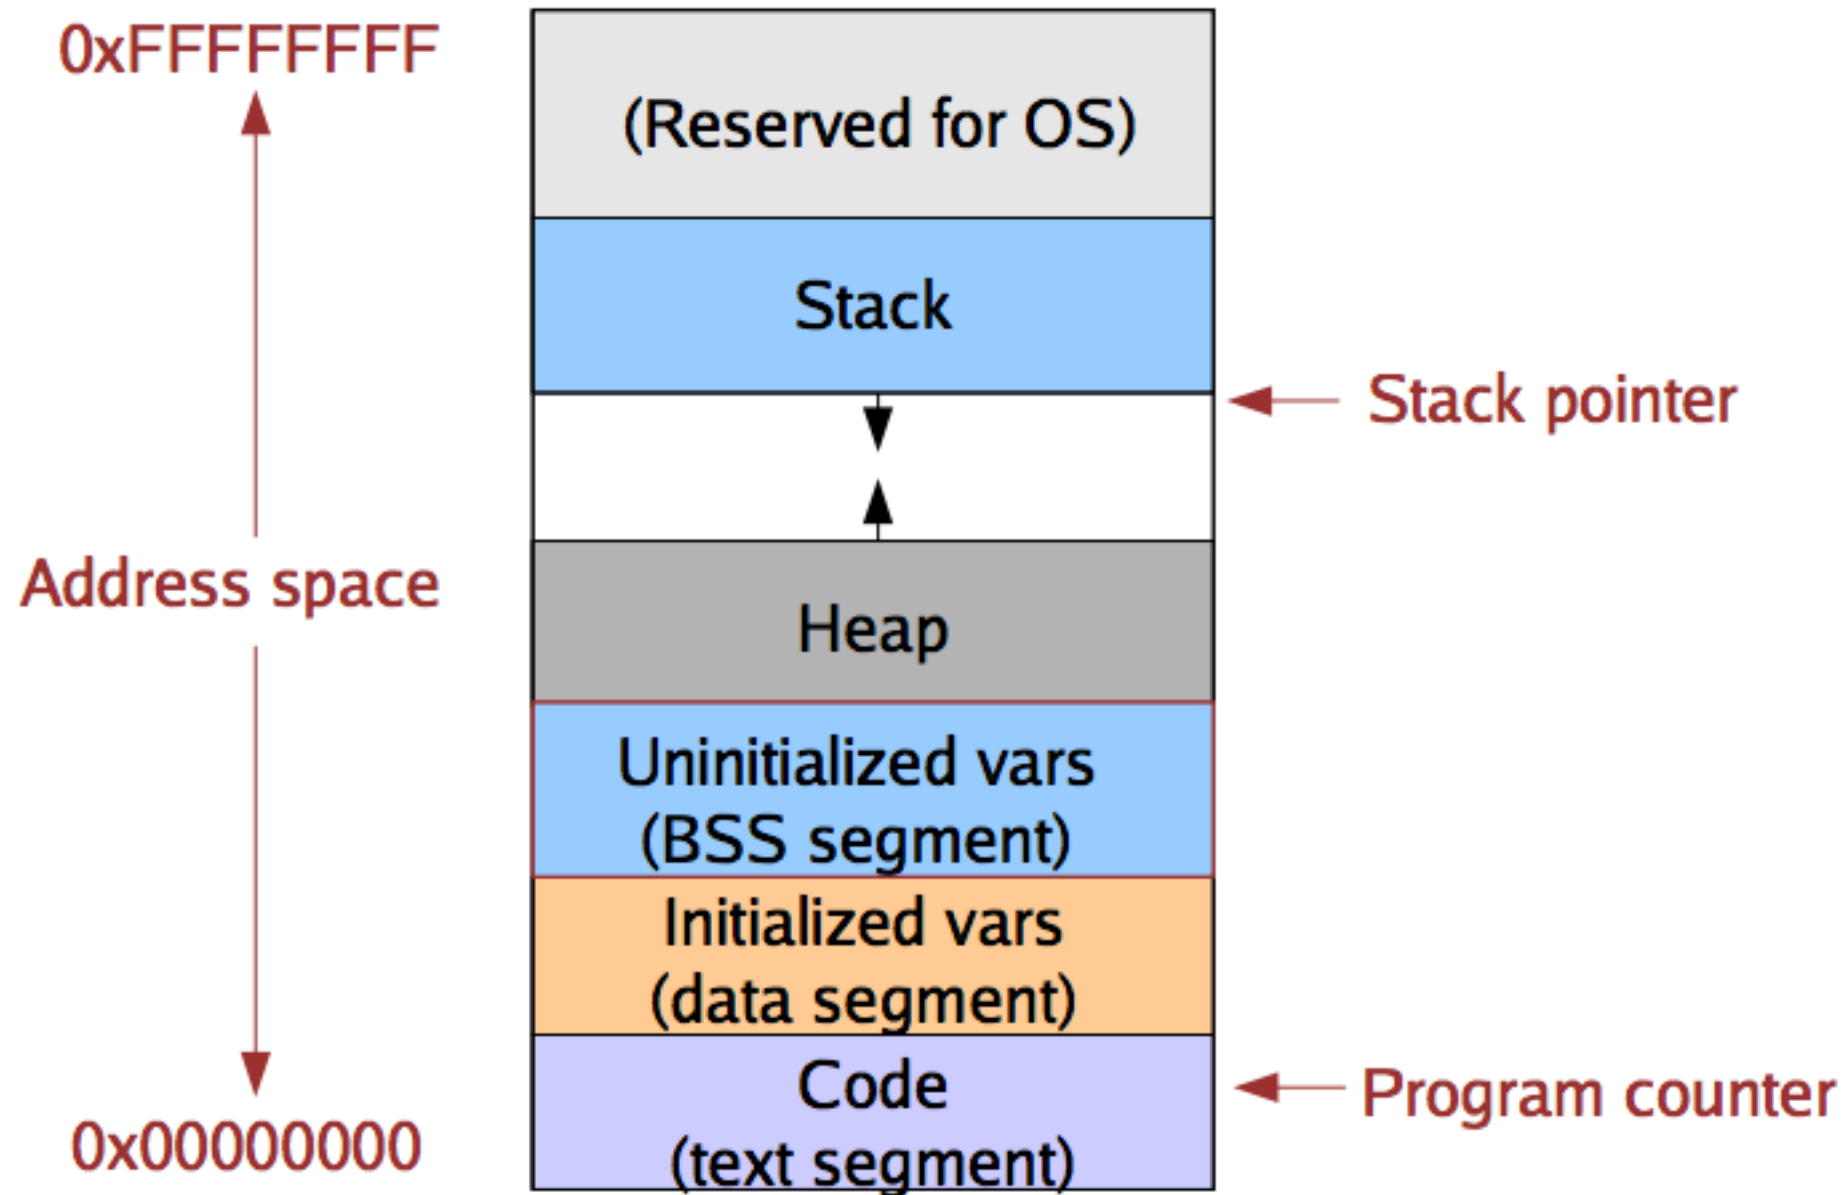
- Implementation

- HW support: address translation
 - Memory management unit (MMU), translation lookaside buffer (TLB)
- OS support: mapping between VA-to-PA
 - Page fault handler, page table

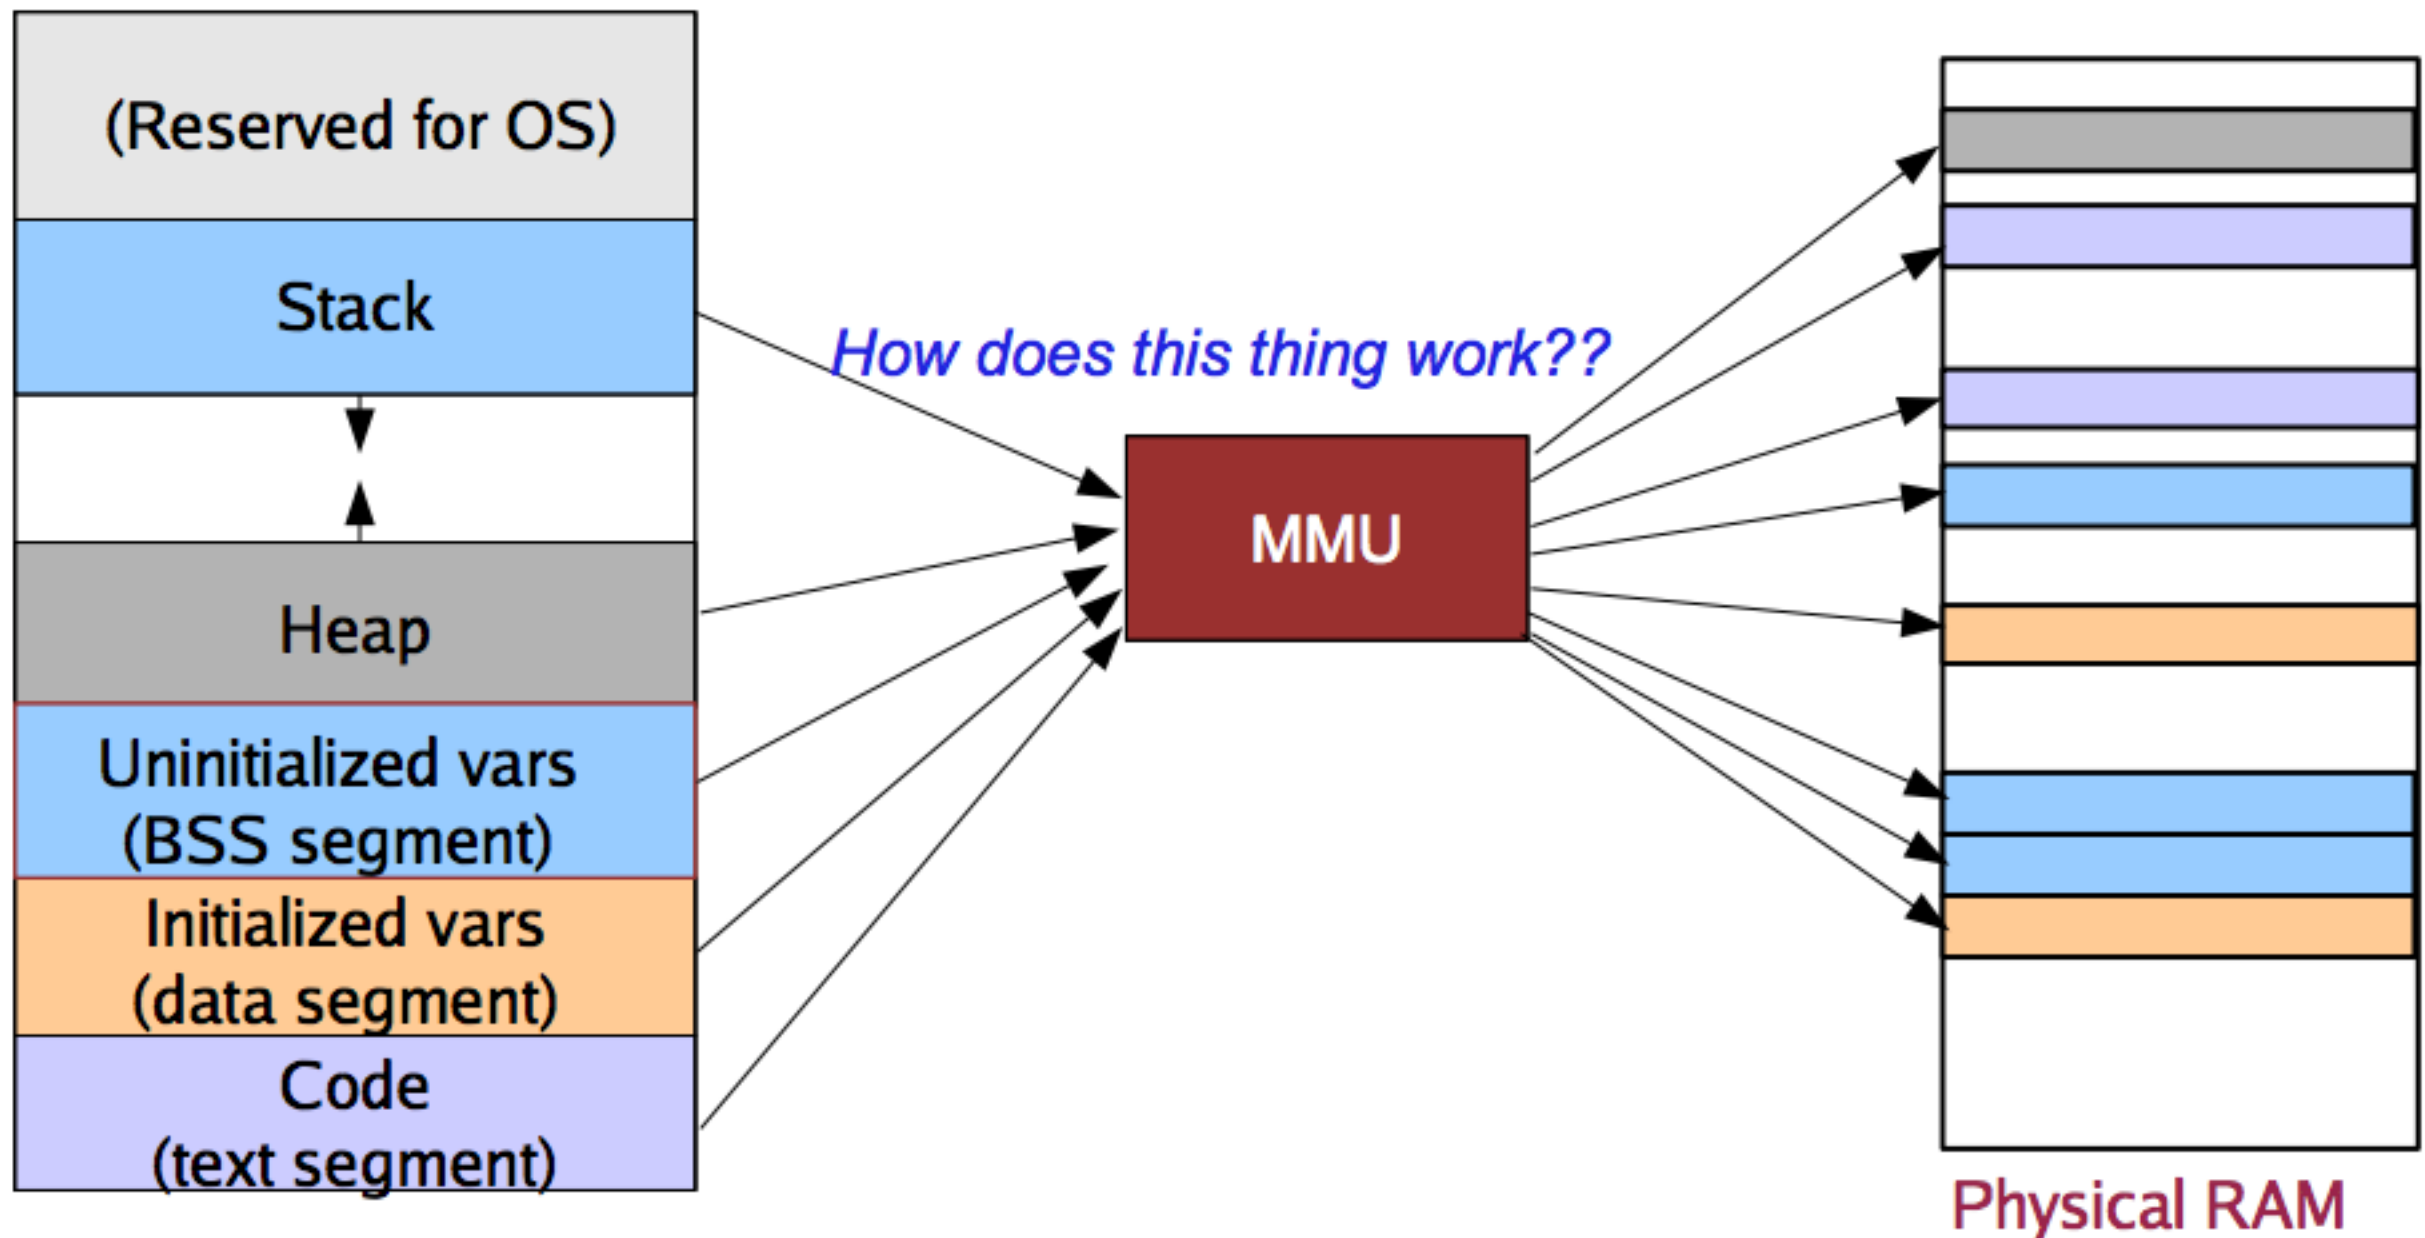
Virtual Address (VA)

- VA: a memory address that a process uses to access its own memory
 - $VA \neq PA$ (physical address)
 - $PA = MMU(VA)$
 - VA-to-PA mapping
 - Determined by OS

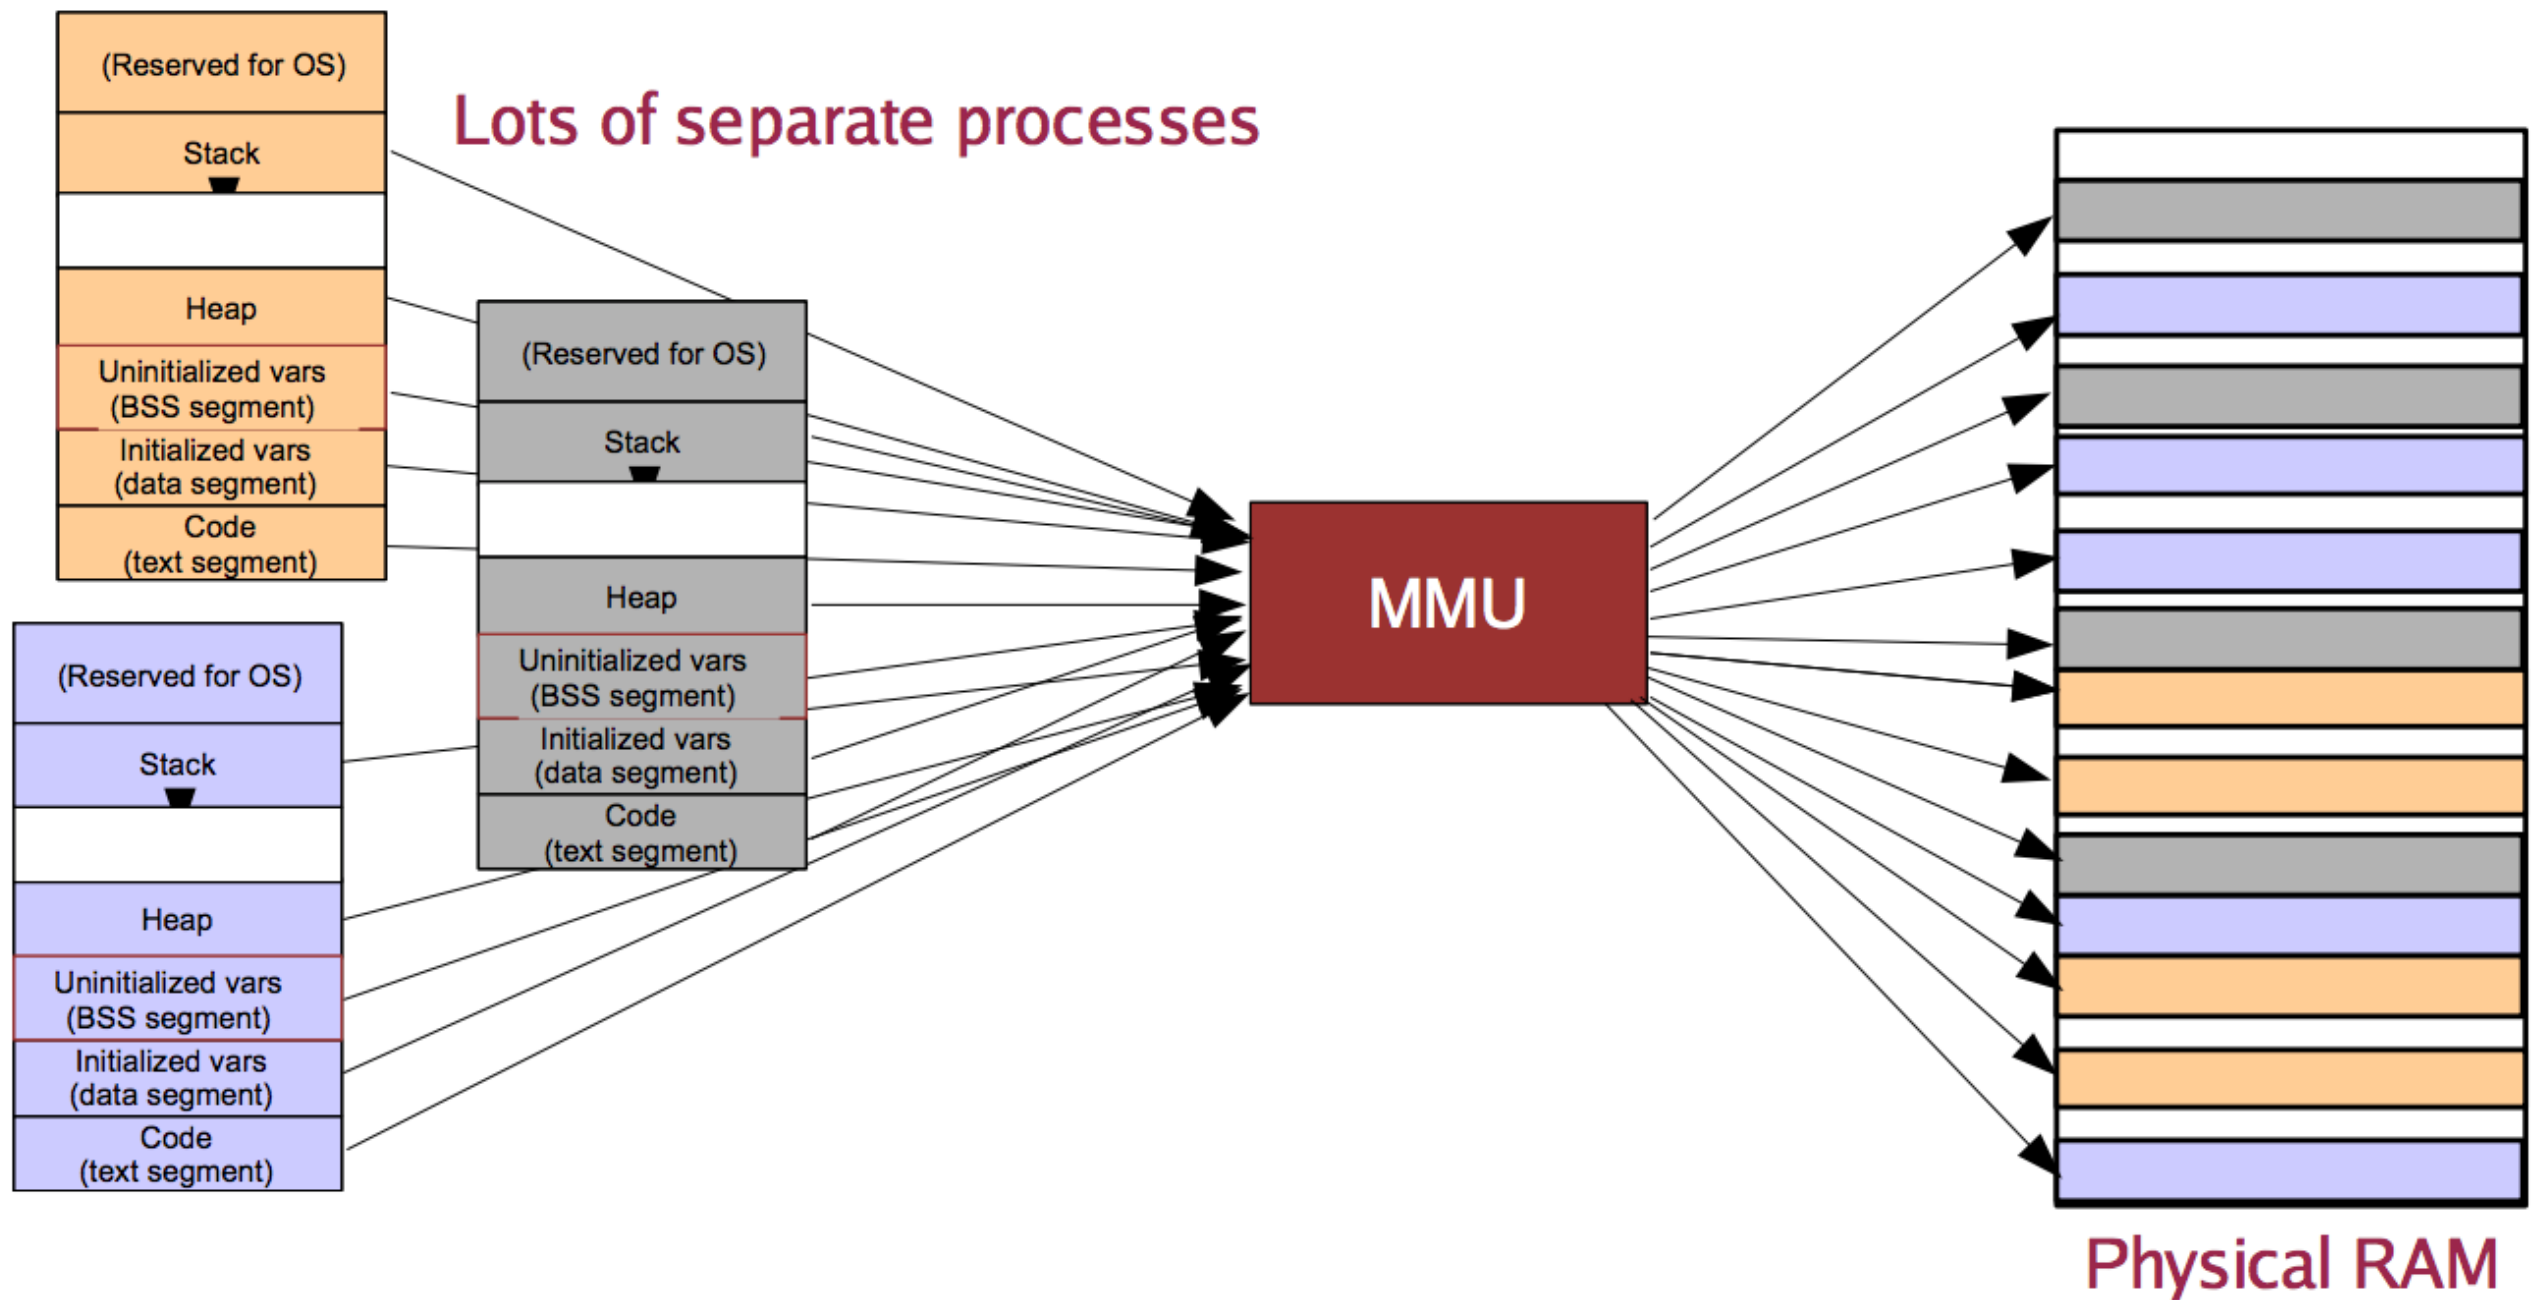
Virtual Address (VA)



Virtual Address (VA)



Application Perspective

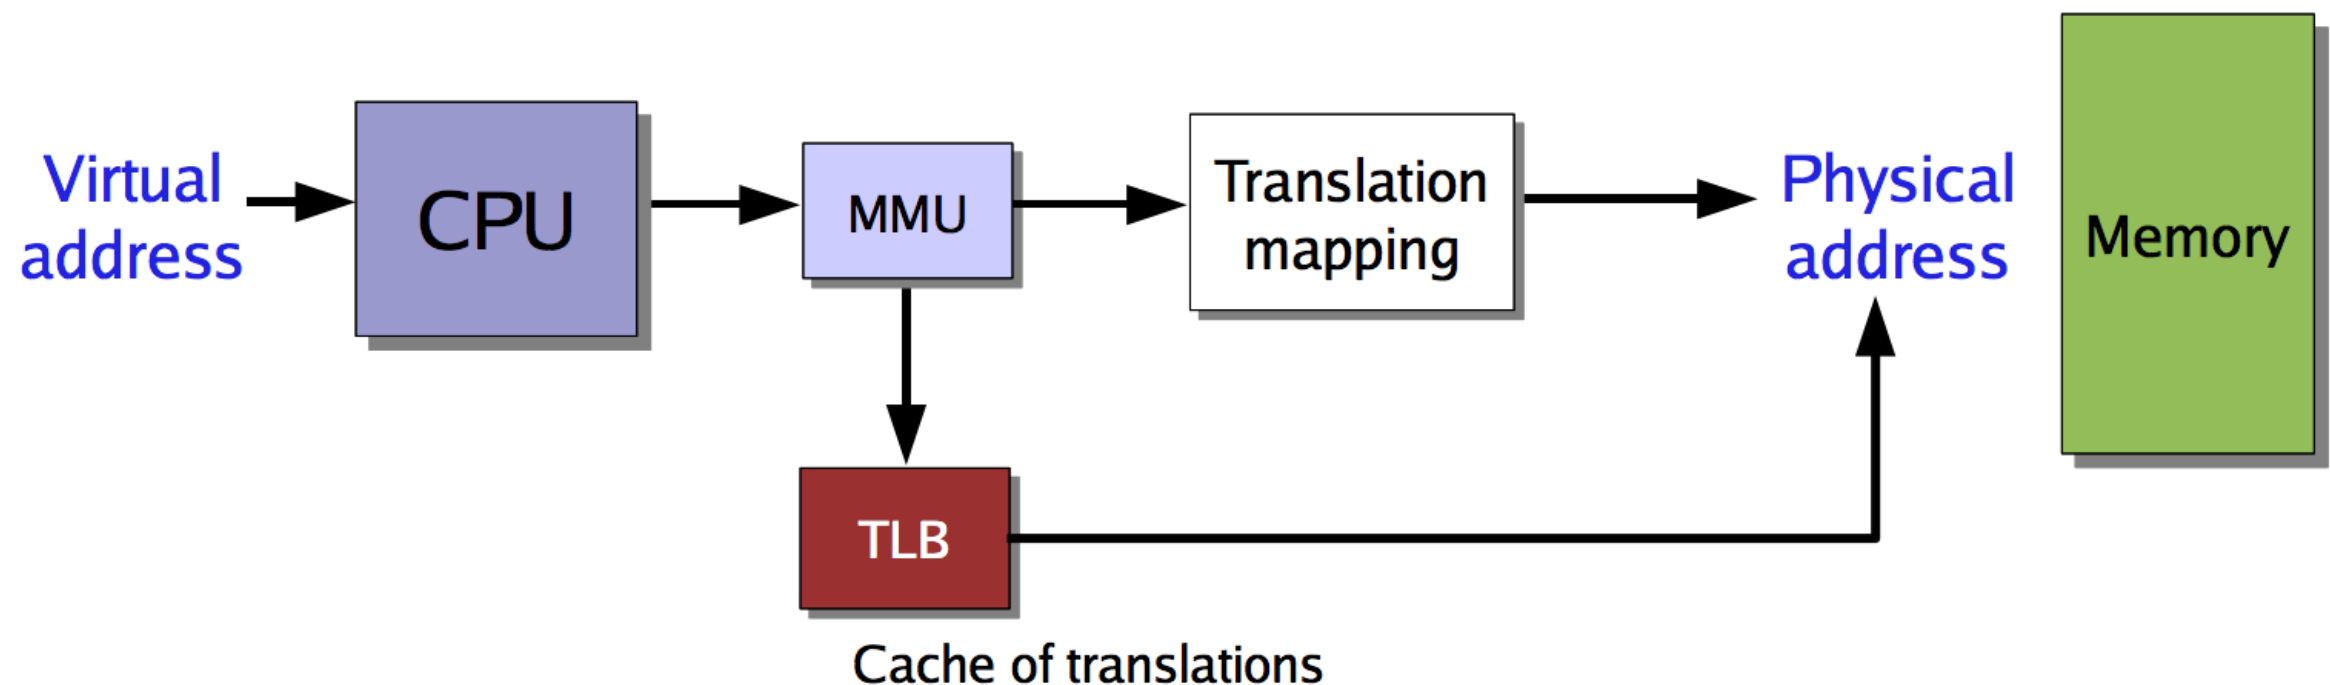


Virtual Address

- Isolation
 - VA in one process refer to different physical memory than virtual addresses in another
 - Exception: shared memory regions btw. processes
- Relocation
 - A program does not need to know which PA it will use when it's run
 - Compilers generate relocatable code

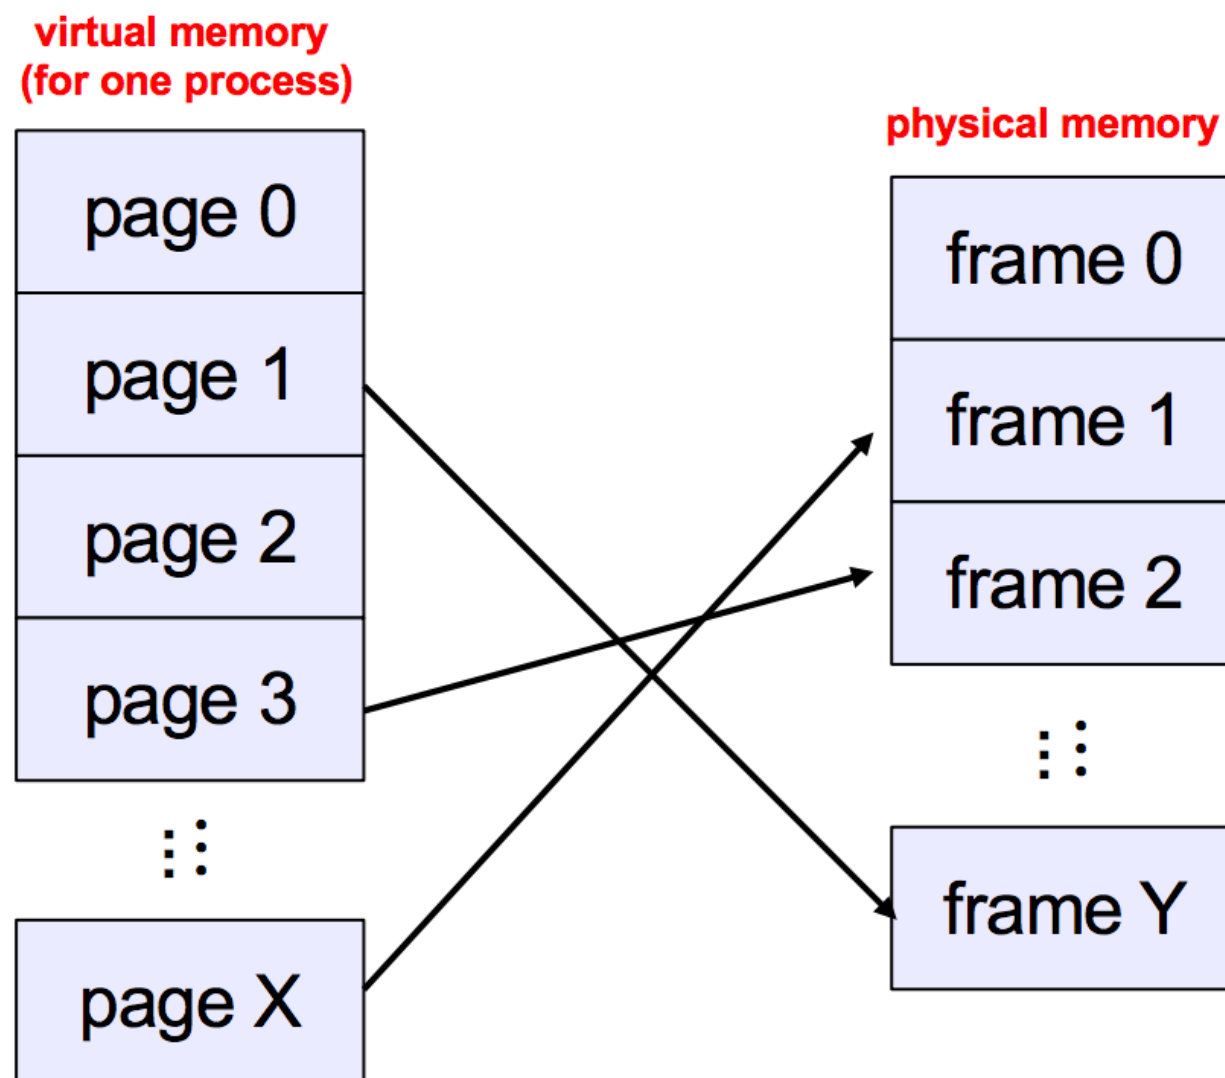
MMU and TLB

- MMU (memory management unit)
 - HW that translates a VA to a PA
 - $PA = MMU(VA)$
- TLB (translation lookaside buffer)
 - Cache for MMU V-to-P address translations



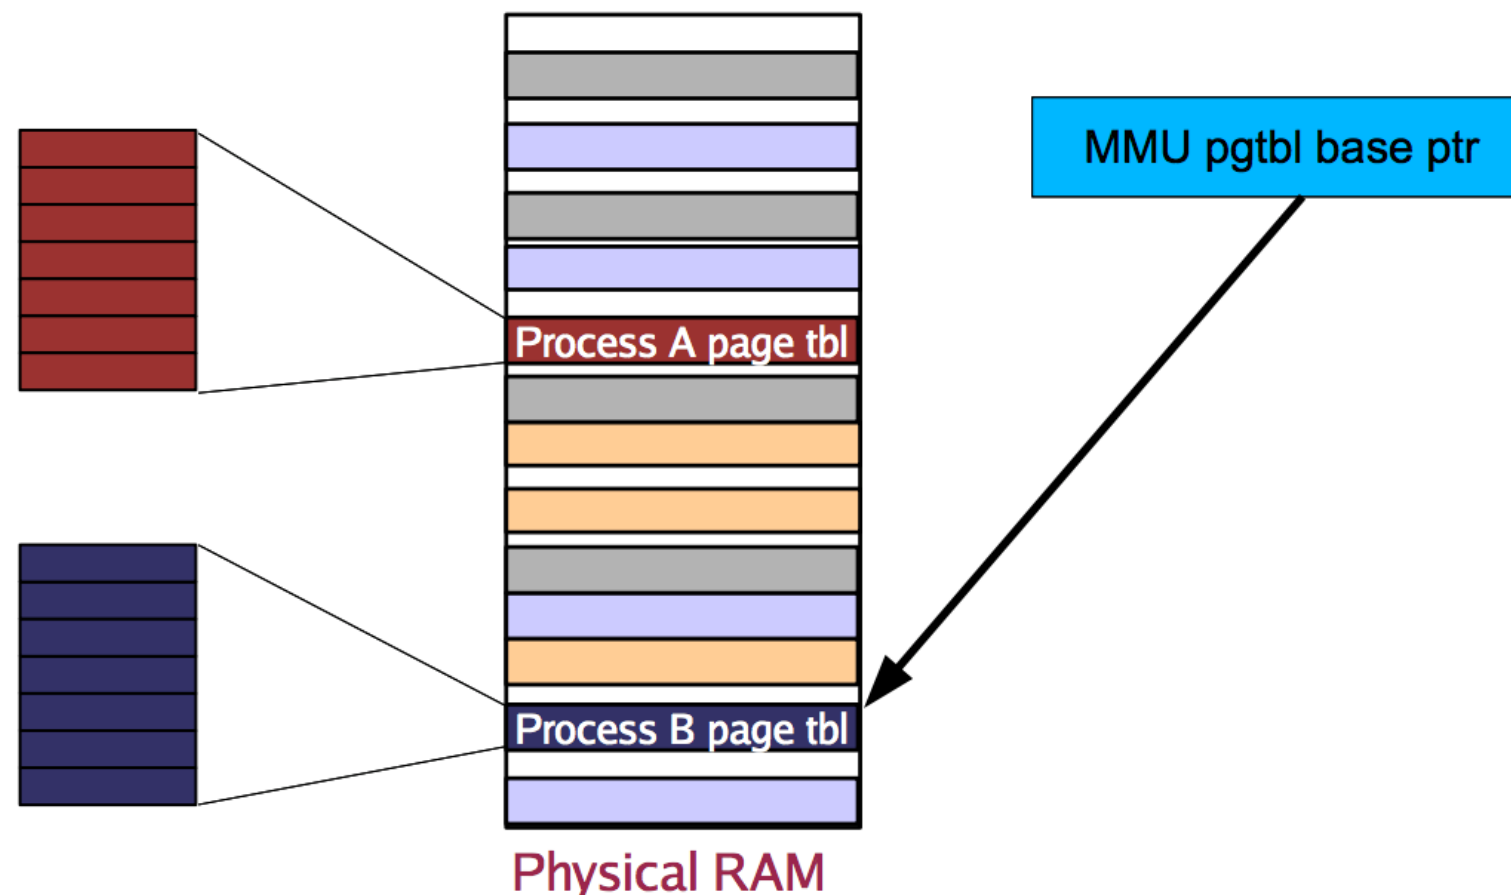
Paging

- Virtual memory is divided into fixed-size chunks called pages
- Physical memory is divided into page frames




Page Tables

- Store the virtual-to-physical address mappings
- Location: in memory
- Way to access
 - MMU has a special register called the page table base register
 - It points to the physical memory address of the top of the page table for the currently-running process



The TLB


- MMU overhead
 - Each memory access requires an additional memory access to the page table (100% overhead)
- Solution: Translation Lookaside buffer (TLB)
 - Very fast (but small) cache directly on the CPU
 - Caches most recent v-to-p address translations
 - A TLB miss requires access to the page table on the main memory



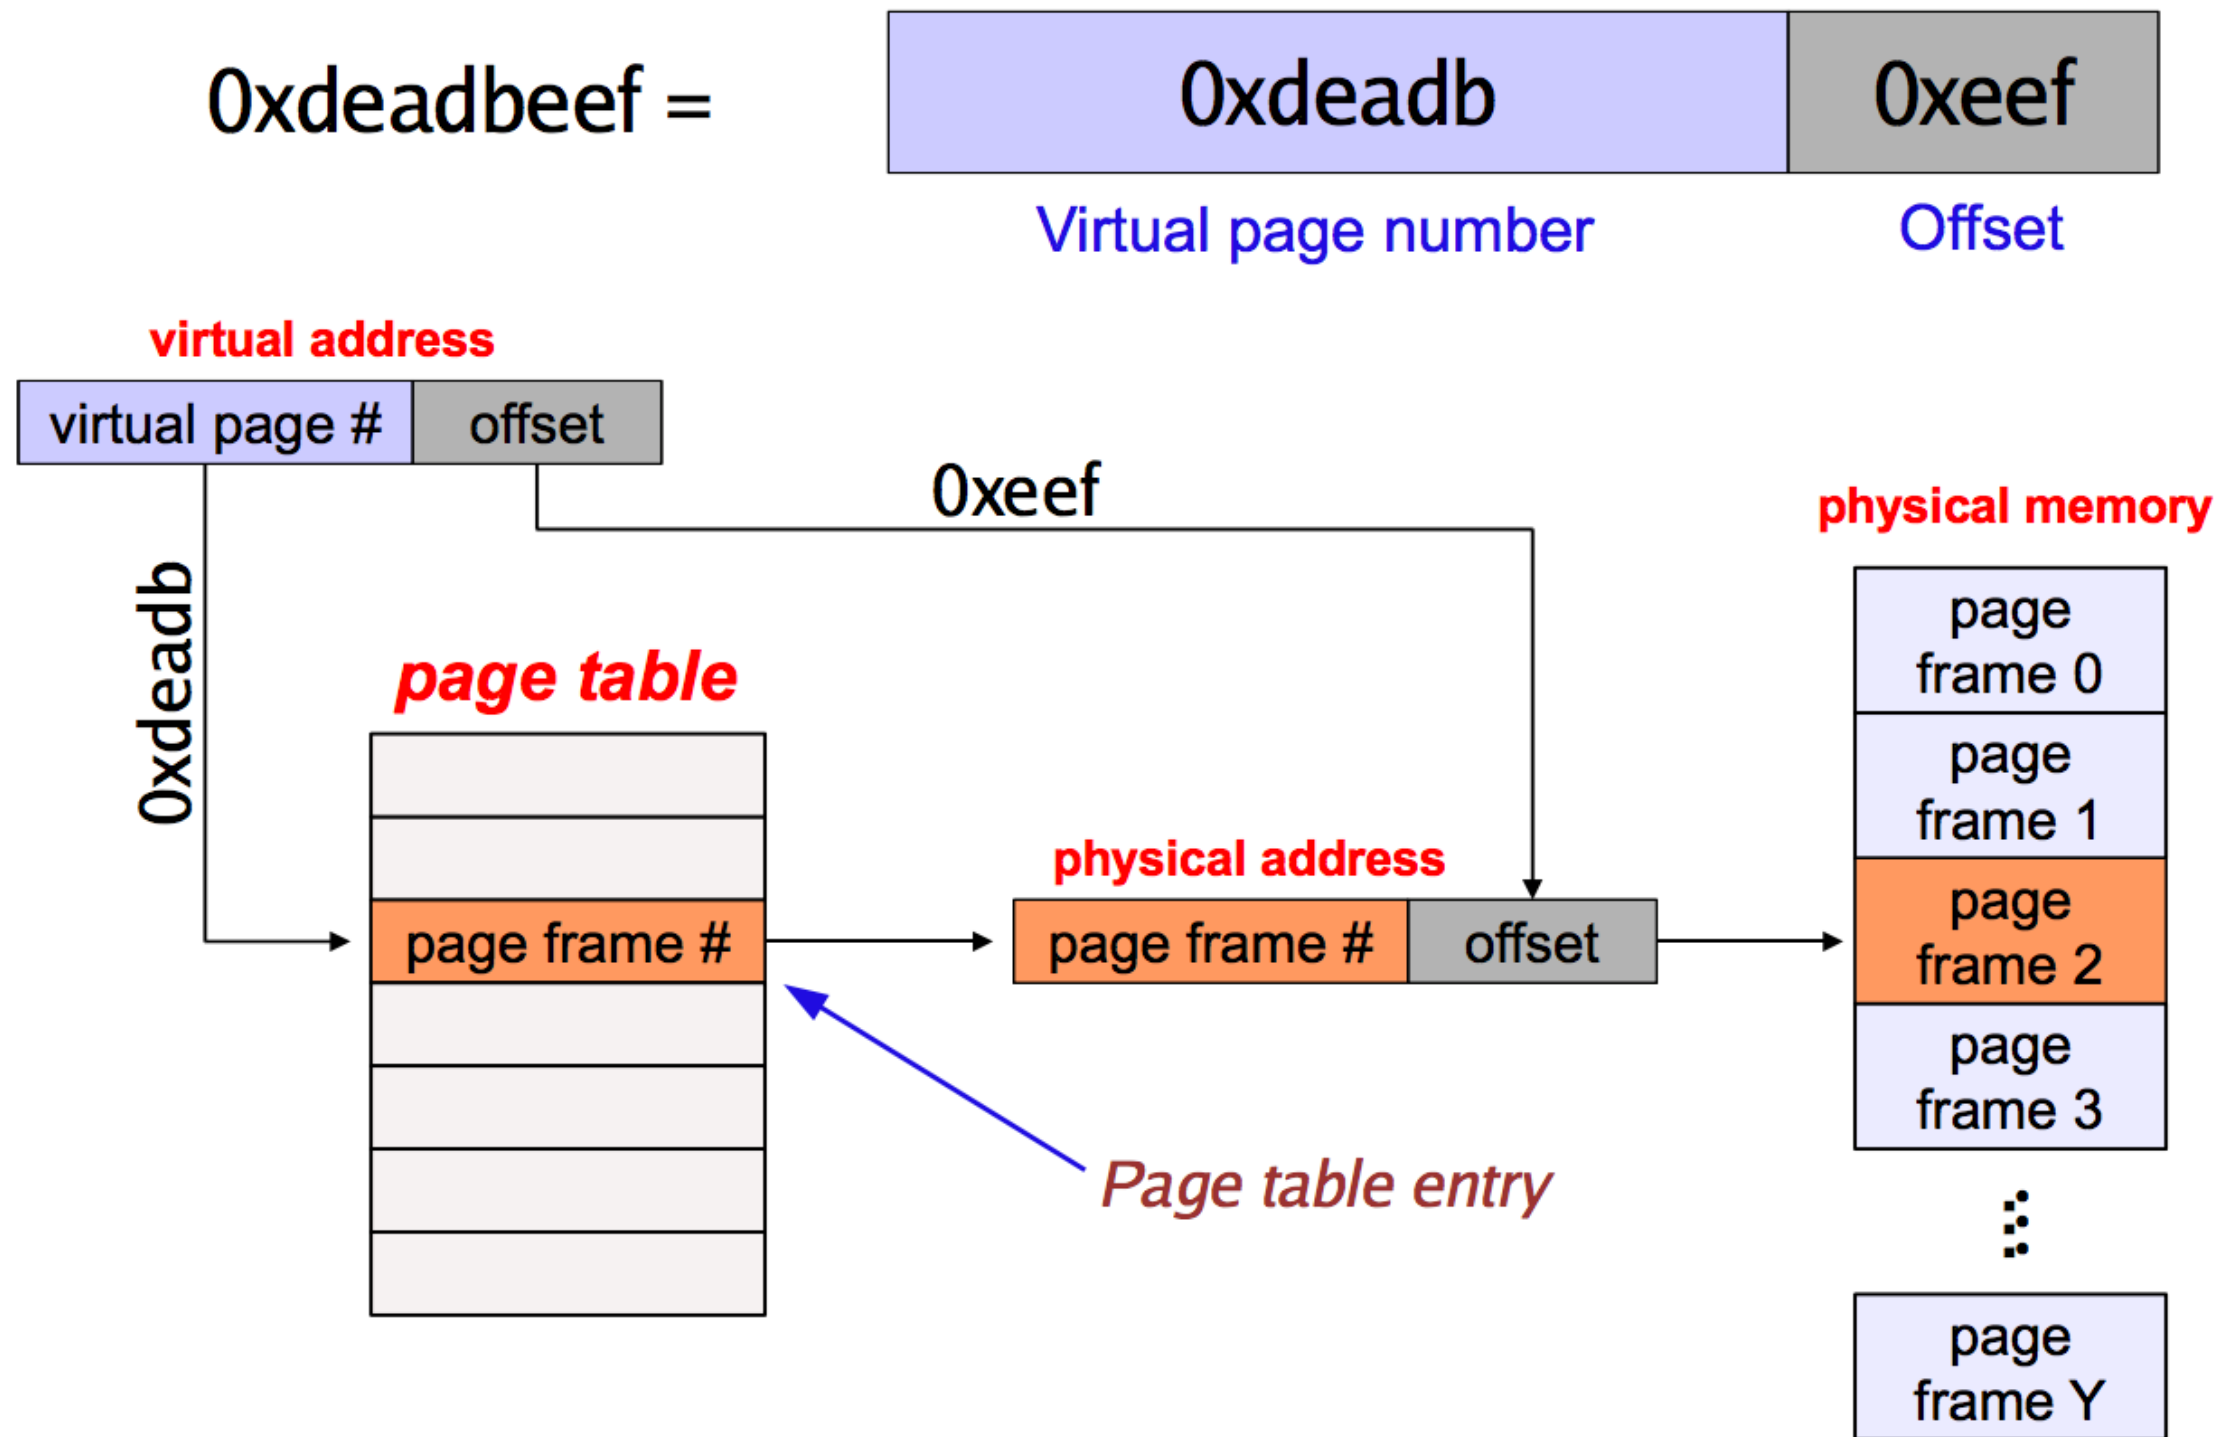
Virtual Address Translation

- Performed by MMU
 - VA is broken into Virtual page number and an offset
 - Mapping from a virtual page to a physical frame provided by a page table

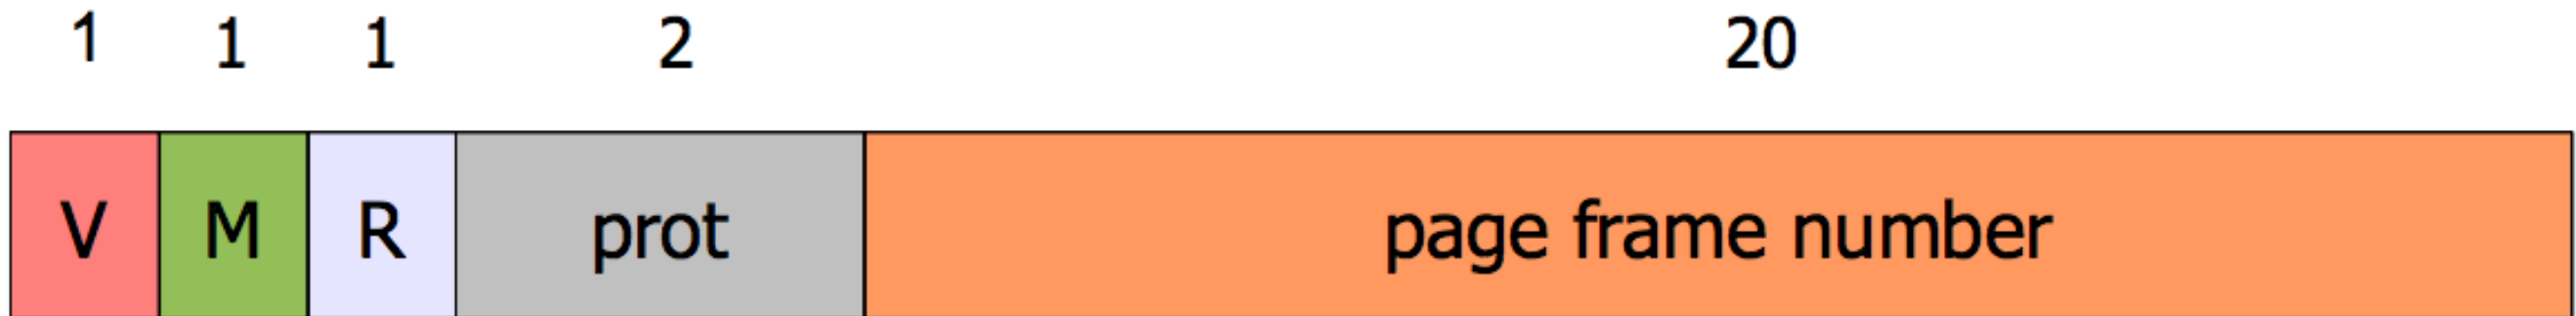
0xdeadbeef =



Virtual Address Translation



Page Table Entries (PTEs)



- Valid bit (V): whether the corresponding page is in memory
- Modify bit (M): indicates whether a page is “dirty”
- Reference bit (R): whether a page has been accessed
- Protection bits: specify if the page is readable, writable, or executable
- Page frame number: physical location of page in main memory
 - $\text{PFN} + \text{OFFSET} = \text{physical memory address}$

Reference

<http://www.eecs.harvard.edu/~mdw/course/cs161/notes/vm.pdf>

https://cs.umd.edu/class/spring2015/cmsc411-0201/lectures/lecture14_virtual_memory_2.pdf

<http://codex.cs.yale.edu/avi/os-book/OS9/slide-dir/PPT-dir/ch9.ppt>